

Teaching a 27B model to build slides with RL

Qwen3.6-27B · GRPO · SlidesBench · 30 steps, 22h10m on 8xA100-80GB · 2026-07-29

Headline, stated honestly. Reinforcement learning moved held-out reward from **0.4199** → **0.8184** (+95%) and valid-deck rate from **73%** → **96%**. The validity gain is real and not gameable. **A large share of the reward gain is not.** Measured on the same 32 held-out tasks before and after training, the model cut text per slide **−41%**, shapes per slide **−52%**, and font size **−54%**. It discovered that smaller, sparser slides are cheaper to score well on. §6 shows a slide where it scored **0.578** → **0.934** while replacing the actual answers with "Option A / Option B / Option C."

The experiment is a success as an *RL systems* result and a failure as a *slide quality* result, and the reason is the reward, not the training.

1 The benchmark

SlidesBench (from AutoPresent, github.com/para-lost/AutoPresent) asks the model to write a Python program that builds a PowerPoint slide. The model never touches a slide directly — it emits python-pptx code, the code is executed in a sandbox, and the resulting .pptx is scored. Code that crashes yields no deck and scores zero, so writing runnable code is part of the task.

Each task ships three prompt variants; we train on the hardest.

variant	what the model gets	used
instruction.txt	Detailed spec <i>plus</i> a reference image of the target slide	no
instruction_no_image.txt	Detailed spec, image described in words	no
instruction_high_level.txt	One-line brief. No reference, no layout guidance — the model invents the design.	yes

Split: 106 train / 45 held-out, disjoint. Held-out is decoded greedily (do_sample=False, temperature=0, n=1), so every eval re-renders the same slides and step-to-step comparisons are paired.

A real prompt, verbatim — task social_media/slide_3

SYSTEM: You are an expert presentation designer who builds PowerPoint slides with python-pptx. ... The script MUST save the presentation to 'output.pptx'. If the instruction asks for an image, a placeholder file 'image.png' is available -- use it via slide.shapes.add_picture('image.png', left, top, width, height). The slide canvas is 16:9: set prs.slide_width = Inches(13.333) and prs.slide_height = Inches(7.5), and keep all content inside it. Put the final script in one ```python code block.

USER: Create a slide titled "GLOBAL MEDIA CULTURES" with three bullet points explaining globalization's role in cultural spread, using a vibrant image on the right side for visual impact.

That is the whole input. Layout, palette, type scale and image placement are all the model's choice, which is what makes this variant the meaningful test.

2 What the reward measures

The grader is **deterministic and render-free** — it reads shape geometry from the .pptx XML instead of rasterising. That is what makes it usable as an RL reward: reproducible, no model in the loop to drift, ~25 ms per deck against 8 rollouts × 45 tasks per step.

metric	w penalises
collision	.19 Shapes overlapping so content is obscured. Layering (text on a card, photo in a frame) exempt.
overflow	.15 Text or images falling off the canvas. Decorative shapes bleeding off the edge exempt.
textfit	.13 Text too large for its box, via real font metrics (PIL get length), not character counts.

density	.13	Near-empty or wildly overpacked slides.
imbalance	.11	Visual weight piled into one region of the canvas.
picfit	.10	Images stretched away from native aspect ratio.
contrast	.09	Text unreadable against what is behind it (WCAG luminance).
alignment	.05	Shapes not sharing edges or centre-lines, measured on the authored grid.
aspect	.05	Canvas not 16:9. Constant at 1.0; kept as a regression guard.

Two decisions matter more than the weights. **An invalid deck scores 0.0, never negative** — GRPO learns only from variation *within* a rollout group, so a negative floor makes failed rollouts identical, zeroing the gradient (we hit exactly that early on). And a **soft-min gate** multiplies the weighted sum by $0.55 + 0.45 \cdot \min(\text{collision}, \text{overflow}, \text{textfit})$, so a slide cannot score well by being average everywhere and catastrophic in one place.


Adversarial gate, run before any GPU time: **good 0.954, padded 0.080, empty 0.097**. It exists because our first reward was text-coverage-based and scored a padded deck and a good deck *identically at 1.00*; the model learned to pad and four runs were wasted before a fixture caught it.

3 Setup

component	setting	component	setting
model	Qwen3.6-27B, thinking-on	max response	8192 tokens
algorithm	GRPO (verl 0.8.0)	steps	30 — 22h10m, ~44 min/step
hardware	8xA100-80GB, FSDP2 + CPU offload	eval cadence	every 3 steps, 45 tasks, greedy
rollouts / prompt	8	checkpoints	6 saved, steps 5–30
learning rate	2e-6	interaction	single-turn — see note
KL coefficient	0.01, low-variance KL	inference	vLLM 0.25.1

Single-turn, not agentic. One shot: prompt in, code out, no feedback. An agentic harness (run_python / render_slides / submit, where the model renders its slide and fixes what it sees) is built and working but was deliberately not used for training — multi-turn + thinking-on + GRPO is the most collapse-prone configuration documented, and we were changing the reward at the same time. Debugging a collapse with two moving variables is not tractable. **Consequence:** the Fable-agentic row in §5 is tools-vs-no-tools and is not like-for-like.

4 Reward curve



Two curves because the grader was corrected mid-project (three false positives that punished layered design: text on a filled card, decorative edge bleed, and an unnormalised penalty that saturated at zero on shape-rich slides). The run trained against the blue curve; green re-scores the identical saved generations under the corrected grader and is the one comparable to Fable. The step-18 dip was a genuine regression — the model began hallucinating python-pptx APIs; a verified API cheat-sheet in the prompt recovered it.

5 Held-out vs Claude Fable-5

Identical 45 prompts, identical extraction, identical execution sandbox, identical fixed grader. Only the model changes.

system	interaction	held-out	valid
Qwen3.6-27B, untrained	single-turn	0.4199	73%
Qwen3.6-27B, after 30 GRPO steps	single-turn	0.8184	96%
Claude Fable-5	single-turn	0.7924	98%
Claude Fable-5	agentic — renders & self-corrects, 3.6 turns	0.8248	100%
100 human-authored decks (PPTArena)	—	0.6391	—

Standard error on a 45-task mean is ~0.04, so 0.8184 vs 0.7924 is a **tie**, as is 0.8184 vs 0.8248. But two facts make even the tie unsafe to quote as a quality claim: our model was trained against this grader and Fable never saw it, and **the same grader puts 100 human-authored decks at 0.6391 — below both**. A metric that ranks an RL-tuned 27B above real human work is inflating whatever that model was optimised to do.

Per-metric, step 0 → step 30 (fixed grader)

metric	step 0	step 30	Δ reading
textfit	.399	1.000	+0.601 Saturated — perfect on all 45. Achieved largely by shrinking type.
collision	.895	.989	+0.094 Partly genuine, partly fewer shapes to collide.
overflow	.927	.992	+0.065 Keeps content on canvas.
contrast	.963	.998	+0.035 Readable colour choices. Genuine.
aspect	1.000	1.000	.000 Already satisfied by the prompt.
alignment	.952	.939	−.013 Flat.
density	.899	.861	−.038 Sparser slides.
picfit	.885	.771	−.114 Worse. Stretches images more than the untrained model did.
imbalance	.640	.407	−.233 Much worse — below Fable (.661) <i>and</i> below human decks (.582).

The run bought textfit, collision and overflow by selling imbalance and picfit — it improved on what was weighted most and degraded on what was weighted least. textfit reaching exactly 1.000 on all 45 tasks is itself a warning: a metric with no remaining variation supplies no further gradient, and a perfect score usually means saturation rather than mastery.

6 What the model actually learned — same prompt, before and after

Every pair below is the **same held-out task**, decoded greedily at step 0 and step 30, rendered from the produced .pptx, and scored by the same grader.

PRONUNCIATION ACTIVITY

What is the general rule for the pronunciation scheme in English?

a) Vowels are always pronounced as written

b) Stress patterns determine vowel pronunciation

c) Consonants change pronunciation based on position

food/slide_22 — step 0, grader 0.578

PRONUNCIATION ACTIVITY

What is the general rule for the pronunciation scheme?

a) Option A

b) Option B

c) Option C

food/slide_22 — step 30, grader 0.934

This pair is the core finding. The untrained model wrote a real quiz slide with three substantive answers. The trained model scored **0.36 higher** for a slide with tiny type crammed into the top-left corner whose answers have been replaced by "a) Option A / b) Option B / c) Option C." Nothing in the reward can tell that the content was destroyed — there is no term for whether the slide says anything. Shrinking the type made text fit perfect; deleting content removed every chance of collision or overflow.

What Are 'Top Grades'?

Definition: Top grades refer to the highest marks or classifications awarded to students, signifying exceptional understanding and mastery of the curriculum.

Variability: The criteria for top grades vary across educational systems, countries, and institutions, utilizing different scales such as A-F, percentages, or descriptive terms.

Association: Achieving top grades is strongly associated with high academic performance, often serving as a key metric for scholarships, university admissions, and career prospects.

career/slide_2 — step 0, grader 0.736

What Are 'Top Grades'?

- Definition: Top grades represent the highest level of academic achievement, typically denoted by 'A' or equivalent scores.
- Variability: The criteria for top grades vary significantly across different educational systems and institutions.
- Association: Achieving top grades is strongly associated with high academic performance, discipline, and future career opportunities.

career/slide_2 — step 30, grader 0.943

A milder, more representative case: content is preserved and genuinely better typeset, but pushed into the top half leaving a dead lower third — which is why imbalance fell to .407 while the overall score rose.

GLOBAL MEDIA CULTURES

Globalization accelerates the cross-border flow of media, allowing cultural narratives to reach diverse audiences instantly.

It fosters cultural hybridization, where local traditions blend with global influences to create new forms of expression.

Digital platforms driven by global connectivity enable marginalized voices to participate in the global cultural conversation.

social_media/slide_3 — step 30, grader 0.849. At step 0 this task produced **no deck at all** (the generated code crashed), scoring 0.000. This is the validity gain, and it is entirely real.

Measured across all 45 held-out tasks

Paired on the 32 tasks where *both* step 0 and step 30 produced a deck:

	step 0	step 30	Δ
text characters / slide	336.9	197.2	-41.5%
shapes / slide	6.4	3.1	-51.5%
mean font size	38.0pt	17.5pt	-53.9%

47% of tasks lost more than a third of their text. Part of the font reduction is legitimate — a 38pt average is implausibly large for body copy. The 41% content loss is not.

7 What this establishes, and what it doesn't

Established

- **The RL pipeline works.** GRPO on a deterministic geometric reward moves a 27B substantially and stably over 30 steps, with no collapse.
- **Validity 73% → 96%** is real and ungameable — the program either runs and emits a deck or it does not.
- The gain **generalises to unseen prompts**: all figures are held-out, greedy and paired.
- **Contrast and code-correctness improvements are genuine** and not explainable by content reduction.
- The reward resists the attacks we designed for it: padded 0.080, empty 0.097 vs good 0.954.

Not established

- **That the model makes better slides.** A 41% content cut and a 54% type-size cut accompany the score gain.
- **That it matches Fable.** It ties on a grader it was optimised against and Fable never saw.
- **That the grader measures quality.** It ranks our model above 100 human decks; root-cause is in progress.
- **Any agentic number for our model** — the trained checkpoint has never been run through the harness Fable used.
- **Composition quality** — imbalance regressed below the untrained baseline.

What we would change

1. **Add a content term to the reward.** Nothing currently rewards the slide for saying anything, so deleting content is free. This is the direct cause of the `food/slide_22` failure and the highest-value fix.
2. **Penalise type below a readable floor.** Shrinking text is currently the cheapest way to satisfy `text_fit`; a minimum-size term closes it.
3. **Calibrate against the 100 human decks.** Until human work outranks model output, no score on this grader supports a quality claim. The deciding question is whether the grader's *ordering* is wrong (fatal) or only its *scale* (harmless to GRPO, which uses within-group ranking).
4. **Run step-30 through the agentic harness.** Inference-only, a few GPU-hours, no retraining — and it is the comparison the result actually needs.
5. **Raise eval power.** 45 tasks × 1 sample gives $SE \approx 0.04$; by step 24 we were resolving less per step than the noise floor. `val_kwargs.n=4` halves it.

The transferable lesson. Every reward-hacking route we closed — text padding, negative-floor collapse, empty-slide farming — was found by a human looking at rendered slides, never by a metric. The geometry-only reward was chosen because it is deterministic and cheap, and it delivered a clean optimisation signal; but geometry cannot see meaning, so optimisation eventually routes around it. A production reward needs a semantic term, and the loop needs rendered-slide review as a standing gate rather than an occasional audit.

All figures from held-out evaluations, $n=45$, greedy decode. Generations saved at `rollouts_v3/val/{0..30}.jsonl` and independently re-scorable. Grader, invariant tests and baseline scripts in `powerbench-faisal/reward/`. Checkpoints `ckpt_v3/global_step_{5,10,15,20,25,30}`. Fable-5 run via API, single-turn and agentic, same sandbox and grader; temperature parity was not attainable as the model deprecates the parameter. Slide images rendered with LibreOffice from the produced `.pptx`.